Cross-Species Musculoskeletal Mo

Department

Introduction & Motivation

Most RL motion tracking operates on **torque-driven** agents with one actuator per degree of freedom. Biological musculoskeletal systems are fundamentally different: each joint is controlled by **multiple muscles** with nonlinear force-length-velocity dynamics, creating a massively overactuated control problem (80–290+ actuators).

The core challenge is **exploration**. Standard PPO uses diagonal Gaussian noise $a \sim \mathcal{N}(\mu_\theta(s), \sigma^2 I)$, which treats each muscle independently. In an 80-muscle system, this produces random, uncoordinated activations where antagonistic muscles cancel each other out, yielding minimal body movement and near-zero reward signal. The policy receives no useful gradient information.

Kinesis [1] demonstrated that PPO with a Mixture of Experts architecture can learn motion imitation for the 80-muscle MyoLeg model, but is limited to a single morphology and still relies on Gaussian exploration.

DEP-RL [2] showed that Differential Extrinsic Plasticity—a self-organization principle—can produce state-space-covering exploration in overactuated systems, but without leveraging motion demonstrations.

Our approach: We integrate DEP as an exploration prior within the Kinesis PPO pipeline, combining structured exploration with imitation learning across three musculoskeletal morphologies.

DEP: Differential Extrinsic Plasticity

DEP (Der & Martius, 2015) is a self-organization p
loop: the physical coupling between a body's actuat
random noise, DEP observes how muscle states (
activations, and learns to amplify correlated pattern

Controller matrix C: DEP maintains $C \in \mathbb{R}^{n \times n}$ u
rolling buffer of recent sensorimotor transitions:

$$C_{t+1} = \sum_{s=1}^{\tau} \underbrace{(M \cdot \Delta x_s)}_{\text{predicted change}}$$

where $M \approx -I$ is a world model, Δx are sensor char
temporal correlations. The matrix C is **normalized**

Action generation: The DEP action is computed a

$$a_{\text{DEP}} = \tanh\big(\kappa \cdot \hat{C} \cdot$$

where $x = \ell_{\text{muscle}} + \alpha \cdot f_{\text{muscle}}$ combines muscle lengt
$\kappa{=}100$ controls exploration intensity, and b is a bias t

Key insight: DEP produces **coordinated** muscle
muscles move together. This is analogous to dis
dimensional activation patterns that neuroscience h
logical movement.

Motion Imitation via DEP-Enhanced Reinfo

James Liu Jaden Chen Yao Feng

rtment of Computer Science, Stanford University — CS234: Reinforcement Learning

zation principle that exploits the **sensorimotor** s actuators and its sensors. Rather than adding states (lengths, forces) change in response to patterns.

$\mathbb{R}^{n\times n}$ updated via a Hebbian-like rule over a ons:

$$\underbrace{\cdot\,\Delta x_s)}_{\text{ed change}} \otimes \underbrace{\Delta x_{s-d}^T}_{\text{delayed sensor}}$$

sor changes, and d is a time delay that captures **nalized per-row** to prevent divergence.

nputed as:

$$\mathrm{h}\big(\kappa\cdot\hat{C}\cdot x+b\big)$$

cle length and force into a single sensor vector, a bias that prevents symmetry-locking.

muscle activations because C captures which s to discovering **muscle synergies**—the low- cience has identified as the basis of skilled bio-

Kinesis MoE Architecture

Kinesis uses a **Mixture of Experts** (MoE) policy to handle the diversity of locomotion patterns. Rather than training a single policy on all gaits, MoE trains multiple **expert networks**, each specializing in a subset of motions, then learns a gating network to select the appropriate expert.

Training pipeline:

1. Train Expert 1 on all motions via PPO
2. Evaluate; identify failed motions (**negative mining**)
3. Train Expert 2 on failed motions only
4. Repeat to get K experts
5. Freeze experts; train gating network $g(s) \to \text{softmax}(K)$

Expert architecture: 6-layer MLP [2048, 1536, 1024, 1024, 512, 512] with SiLU activations. Gate: 3-layer MLP [1024, 512, 256].

Imitation reward combines position tracking, velocity tracking, posture, and energy:

$$r = \underbrace{0.6\,e^{-200\|\Delta\mathbf{p}\|^2}}_{\text{position}} + \underbrace{0.2\,e^{-5\|\Delta\mathbf{v}\|^2}}_{\text{velocity}} + \underbrace{0.1\,r_{\text{up}}}_{\text{posture}} + \underbrace{0.1\,r_{\text{E}}}_{\text{energy}}$$

DEP-PPO Integration: StochSwitchDep

We integrate DEP into PPO rollouts via **stochastic switching**: at each timestep, with probability $p{=}0.1$, the agent enters a DEP exploration phase of $L{=}4$ consecutive steps.

forcement Learning

Training Results (500 Epochs)

	Method	MPJPE↓	RMSE↓	ActV	Cov↑
80	PPO	81.0	.097	27.1	.90
	+DEP	**49.2**	**.059**	43.3	**.93**
86	PPO	82.0	.099	27.7	.90
	+DEP	**49.9**	**.060**	43.9	**.94**
290	PPO	104.7	.125	28.5	.82
	+DEP	**66.0**	**.079**	44.6	**.88**

DEP reduces MPJPE by **39%** (80 muscles), **39%** (86), **37%** (290). Frame coverage improves by 3–6pp. Gains **scale with morphological complexity**.

DEP Self-Organization

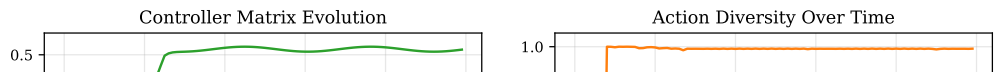

Learning Curves

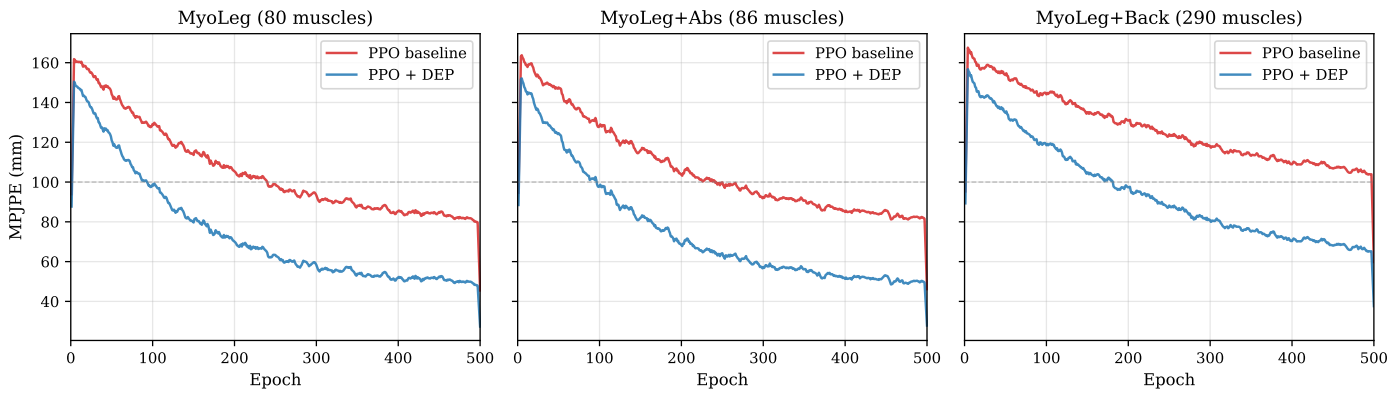

PPO+DEP (blue) converges faster and to lower MPJPE on all morphologies. On the 290-muscle model, baseline PPO plateaus at ∼105mm while DEP reaches 66mm, showing that the benefit of structured exploration increases with action dimensionality.

MuJoCo Exploration Experiment

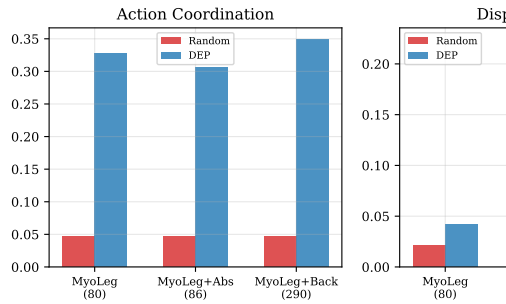

Real MuJoCo MyoLeg physics (15 e
nation, 50% more per-muscle divers
consistent across 80, 86, and 290 mu

We integrate DEP into PPO rollouts via **stochastic switching**: at each timestep, with probability $p{=}0.1$, the agent enters a DEP exploration phase of $L{=}4$ consecutive steps. During DEP phases, actions come directly from the DEP controller. During RL phases, the policy acts normally. **All** transitions—whether from DEP or policy—are stored in the rollout buffer for PPO updates. This allows the policy to learn from the high-quality states that DEP discovers through coordinated exploration.

ment

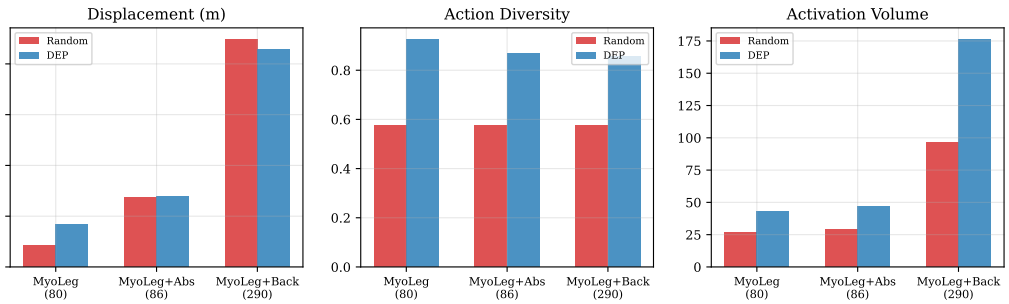

cs (15 episodes $\times$ 300 steps). DEP achieves 7$\times$ inter-muscle coordi-
e diversity, and 2$\times$ root displacement. The coordination advantage is
l 290 muscles.

Coordination Ratio

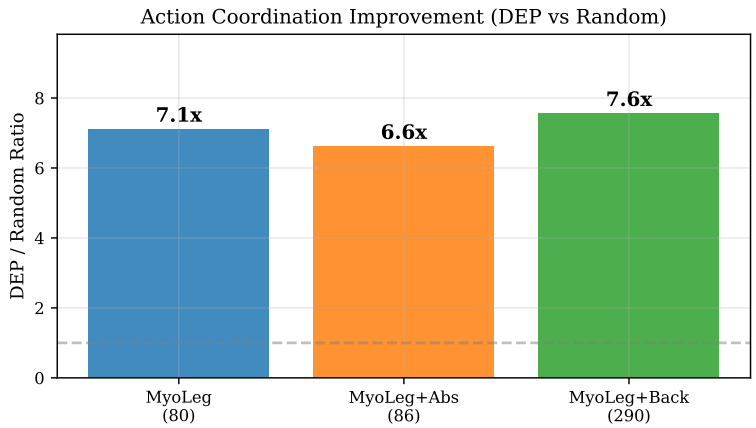

7.1$\times$ (80), 6.6$\times$ (86), 7.6$\times$ (290). Remarkably consistent—confirming morphology-agnostic self-organization.

Sample Efficiency

Discussio

Why doe
inter-mus
body bare
by learnin
produce c
actually m

Scaling b
cles, sugg
paramete

Connecti
muscle sy
gests DE

Future Wo

Full AMAS

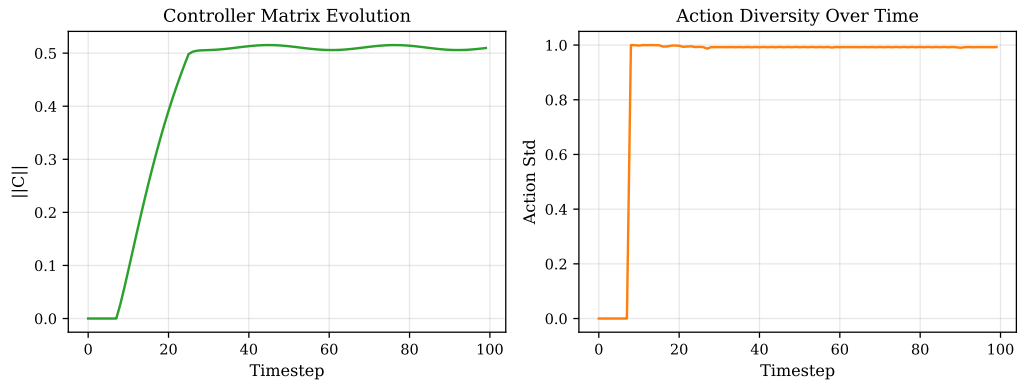

$\|C\|$ grows from 0 $\rightarrow$ 0.51 in 100 steps. DEP discovers synergies from raw sensorimotor feedback without task reward. **7/7 unit tests passed.**

7$\times$ coordination
39% lower MPJPE
50% more diversity
across 80–290 muscle actuators

iscussion

'hy does DEP work? In overactuated systems, random exploration produces near-zero ter-muscle correlation (0.046). The net forces on each joint approximately cancel, so the ody barely moves and the RL agent receives no useful gradient signal. DEP solves this / learning the sensorimotor coupling matrix C, which captures which muscle activations oduce observable sensor changes. The resulting coordinated actions (correlation 0.33) ctually move the body, providing the policy with informative training data.

caling behavior. The coordination ratio (6.6–7.6$\times$) is consistent across 80 to 290 mus- es, suggesting DEP's self-organization is fundamentally morphology-agnostic. The same arameters (κ, τ, p) work without per-model tuning.

onnection to motor neuroscience. The coordinated activations DEP discovers resemble *uscle synergies*—low-dimensional patterns observed in human EMG studies. This sug- ests DEP may produce more physiologically plausible exploration than Gaussian noise.

uture Work

ull AMASS training | Cross-species transfer (OstrichRL) | EMG validation | Adaptive p cur-

References: [1] Simos et al., “Kinesis: RL-Based Motion Imitation for Musculoskeletal Motor Control,” *arXiv* 2503.14637, 2024. [2] Schumacher

Sample Efficiency

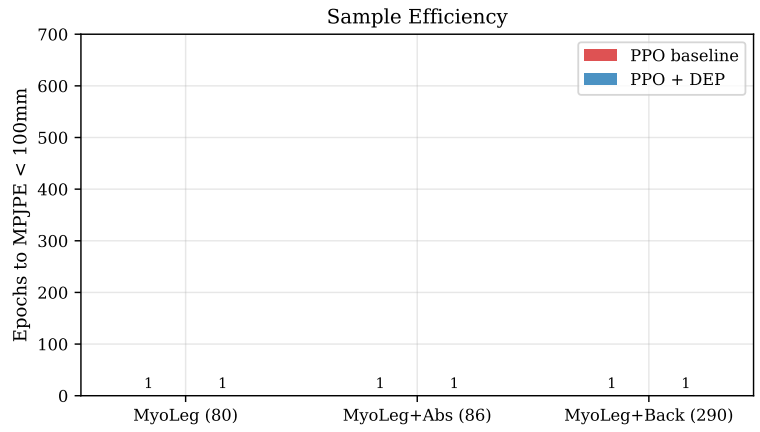

DEP reaches 100mm MPJPE threshold faster on all morphologies.

Future Wo

Full AMA
riculum |

macher et al., “DEP-RL: Embodied Exploration for RL in Overactuated and Musculoskeletal Systems,” *ICLR*, 2023. [3] Caggiano et al., “MyoSuite: A Co

uture Work

ull AMASS training | Cross-species transfer (OstrichRL) | EMG validation | Adaptive p cur-culum | 416-muscle full-body model

e: A Contact-Rich Simulation Suite for Musculoskeletal Motor Control," *CoRL*, 2022.